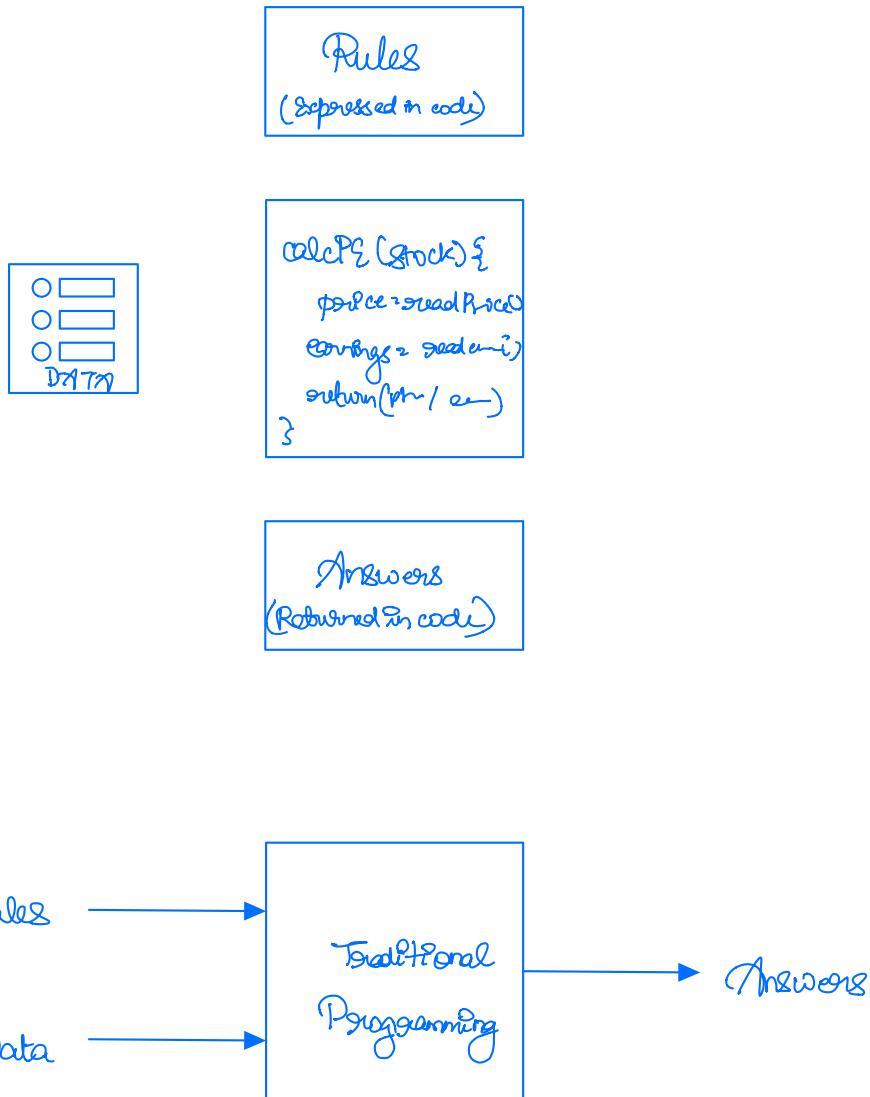
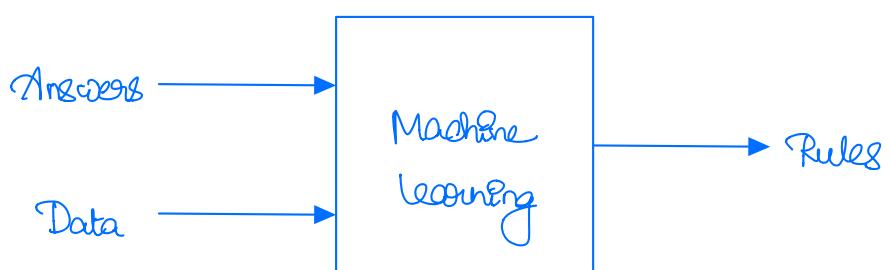


Traditional Programming



Machine Learning



Example : Activity Recognition

In traditional programming, you can tell if a person is walking by checking if their speed is less than 4, running if speed is > 4 but < 12 , and cycling if more than 12.

But what in case of golfing? Similarly, the speed of a cycle can depend upon inclination and surface friction, and people walk at different paces.



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```

// Uh oh

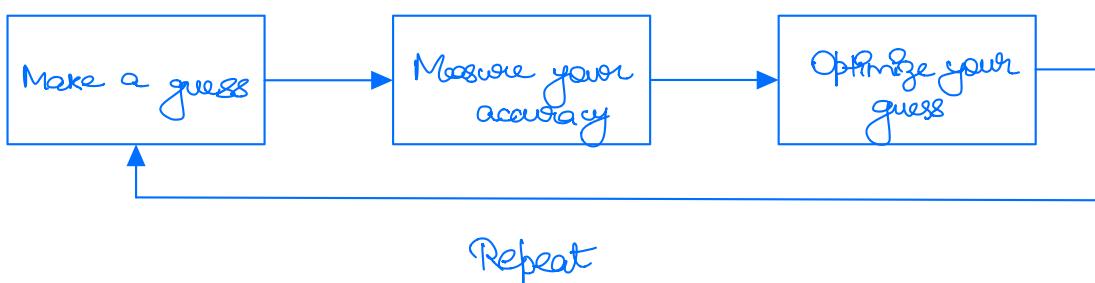
What if we give our data to Machine learning?

It could look like this:

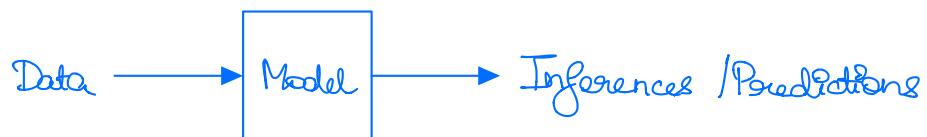
			
010100101010010101 010010101010010111 010100101010010101 001010100101010010 1010 Label = WALKING	101010010100101010 101010100100100100 01001001111010101 11110101110101011 101010101111010101 010101111000111101 1011 Label = RUNNING	100101001111101010 111010101110101011 101010101111010101 01111111000111101 0101 Label = BIKING	111111111101001110 100111110101111101 010101110101010101 110101010101010011 1110 Label = GOLFING (Sort of)

The ML algorithm can look for patterns in the data and assign them to labels, and predict future labels through those learnt patterns, and this can be more suitable than writing rules manually.

The Machine Learning Paradigm

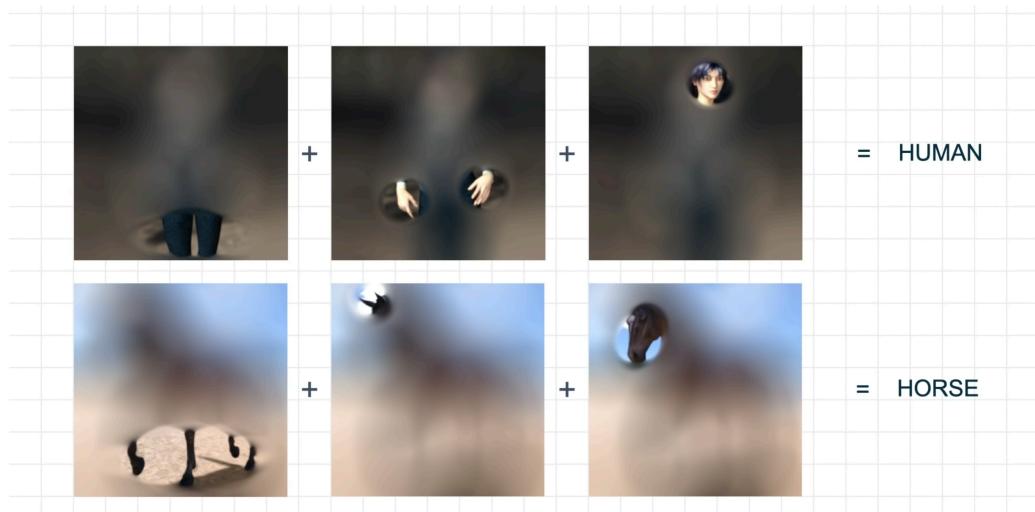


When we do this, we get a model



Computer Vision

Images are just a bunch of numbers. We can add filters to images that a computer can detect that helps a computer understand the image.



If the computer can figure out which filters help extract features that determine between answer of human or horse, then it can differentiate and in future predict human or horse.

In summary,:

- Filters extract features like hands or ears.
- Filters can then be combined with labels to make a prediction of image contents.
- Filters that match the labels are learnt over time. These filters are randomly initialized but over time, using back propagation, their values can be measured for accuracy and optimized.

In terms of code:

In tensorflow, we can define a model as a number of layers, where each layer is a piece of functionality that defines what we want a machine to learn.

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu',
                          input_shape=(300,300,3)),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid'),
])
```

Conv2D stands for 2D Convolution, which is another word for a filter. Since this is a sophisticated scenario with detailed images, we have to apply filters multiple times.

Max Pooling is a way of compressing the image while enhancing features. It is optional but highly useful, as it means that the next steps have less pixels to work with, while still having features.

After all of the filters are learnt, then we use the dense neural network that matches filters to labels, so human legs get matched with humans, and horse legs with horse, and so on.

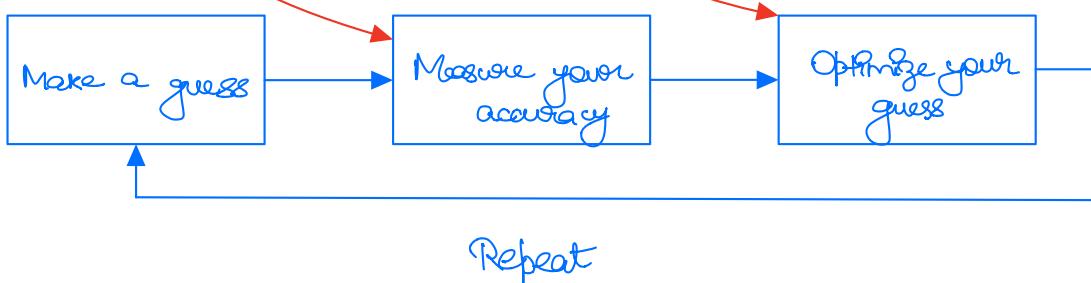
The second dense is the output of the neural network that is also a neuron. As we only have 2 classes here (horse & human), we can do it with a single neuron here, where if the value is closer to 0, then it's a horse, and if value is closer to 1, then it is a human.

Once a model is defined, we have to define how it works with training.

```
optimiser = tf.keras.optimizers.RMSprop(lr=0.001)  
model.compile(loss='binary_crossentropy',  
              optimizer=optimizer,  
              metrics=['accuracy'])
```

① loss function (using Binary Crossentropy as there are 2 classes)

RMSprop stands for Root Mean square propagation and is an optimizer.

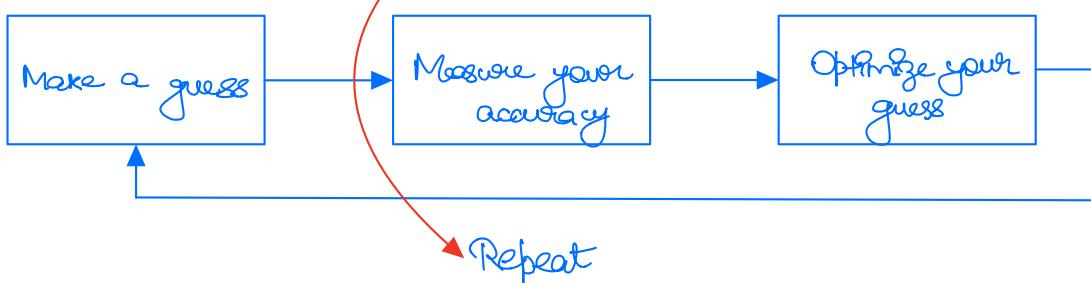


Optimizer takes in the (randomly initialised) parameters accuracy, and update the parameters with a new guess.

Next step is to fit the answers to the labels. In order to figure out the scales:

```
model.fit(train_generator, epochs=12, ...)
```

Defines how many times to repeat the process.



As the model gets better at matching the data to the label, the process convergence is used.

It takes time to figure out how many epochs, what model architecture, what parameters etc. will lead to convergence. A skilled machine learning developer should be able to reach convergence quickly.